

WHAT ELSE HAS MY COMPILER DONE FOR ME LATELY?

UNBOLTING THE COMPILER'S LID...AGAIN

Matt Godbolt
@mattgodbolt
matt@godbolt.org

OUTLINE

- Compiler Explorer story
- Assembly 101
- What else has my compiler done for me lately?
- Behind the scenes of Compiler Explorer

BACKSTORY

```
int sum(const vector<int> &v) {  
    int result = 0;  
    for (size_t i = 0; i < v.size(); ++i)  
        result += v[i];  
    return result;  
}
```

```
int sum(const vector<int> &v) {  
    int result = 0;  
    for (int x : v) result += x;  
    return result;  
}
```

Is one better than the other?
Look at compiler output!

WARNING

- Reading assembly alone can be misleading
- *Always measure too*
- Google Benchmark
- quick-bench.com

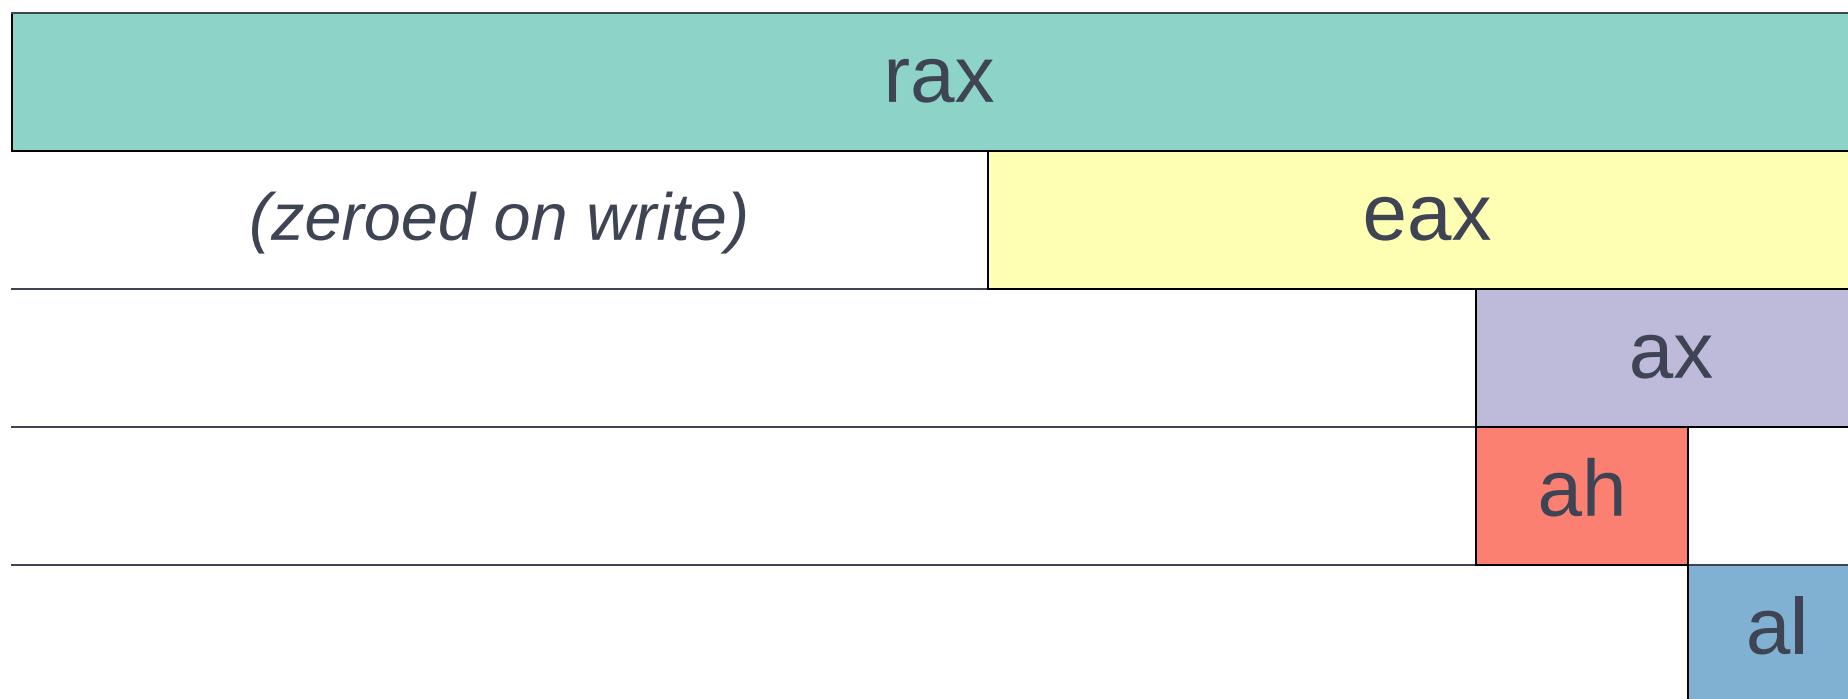
X86 ASSEMBLY 101

REGISTERS

- `rax, rbx, rcx, rdx, rsp, rbp, rsi, rdi, r8-r15`
- `xmm0-xmm15`
- `rdi, rsi, rdx...` arguments
- `rax` is return value

REGISTERS

63...56 55...48 47...40 39...32 31...24 23...16 15...8 7...0



And so on for rdi, rsi, rcx, rdx...

INSTRUCTIONS

```
op  
op dest  
op dest, src  
op dest, src1, src2
```

- op is e.g. call, ret, add, sub, cmp...
- dest, src is register or memory reference:
 $[base + reg1_{opt} + reg2 * (1, 2, 4 \text{ or } 8)_{opt}]$

(Intel asm syntax)

INSTRUCTIONS

```
mov eax, DWORD PTR [r14]
add rax, rdi
add eax, DWORD PTR [r14+4]
sub eax, DWORD PTR [r14+4*rbx]
lea rax, [r14+4*rbx]
xor edx, edx
```

```
int eax = *r14;    // int *r14;
rax += rdi;
eax += r14[1];
eax -= r14[rbx];
int *rax = &r14[rbx];
edx = 0;
```

SUMMARY

- Registers: `rax`, `rbx`, `rcx` ...
- Size: `rax`, `eax`, `ax` ...
- Params: `rdi`, `rsi`, `rdx`, `rcx` ...
- Result: `rax`
- `op dest, src`
- `dest, src` are registers or memory

WHERE WERE WE?

```
int sum(const vector<int> &v) {  
    int result = 0;  
    for (size_t i = 0; i < v.size(); ++i)  
        result += v[i];  
    return result;  
}
```

```
int sum(const vector<int> &v) {  
    int result = 0;  
    for (int x : v) result += x;  
    return result;  
}
```

Which is better?

COMPILER EXPLORER V0.1

```
$ g++ /tmp/test.cc -O2 -c -S -o - -masm=intel \
| c++filt \
| grep -vE '\s+\.'
```

```
sum(std::vector<int, std::allocator<int> > const&):
.LFB786:
    mov rcx, QWORD PTR [rdi]
    mov rax, QWORD PTR 8[rdi]
    sub rax, rcx
    shr rax, 2
    mov rsi, rax
    ...
```

COMPILER EXPLORER V0.1

Not very pretty
TO THE WEB!



DEMO

The screenshot shows a debugger interface with two panes. The left pane displays the C++ source code:

```
1 // setup...
5 int sum(const vector<int>
6     int result = 0;
7     for (size_t i = 0; i <
8         result += v[i];
9     return result;
```

The right pane shows the generated assembly code for an x86-64 architecture using gcc 8.1:

```
1 sum(std::vector<int>
2     mov rdx, QWORD PTR
3     mov rcx, QWORD PTR
4     sub rcx, rdx
```

Both panes have dropdown menus at the top, including "Save/Load", "C++", "x86-64 gcc 8.1", and "O2 -std=c++1z -march=haswell". The assembly code pane also includes tabs for "11010", ".LX0:", ".text", "//", "\s+", "Intel", "Demangle", "A-", and "Libraries".

WALKTHROUGH

```
int sum(const vector<int> &v) {
```

```
; rdi = const vector<int> *
mov rdx, QWORD PTR [rdi]    ; rdx = *rdi      ≡ begin()
mov rcx, QWORD PTR [rdi+8]  ; rcx = *(rdi+8) ≡ end()
```

```
template<typename T> struct _Vector_impl {
    T *_M_start;
    T *_M_finish;
    T *_M_end_of_storage;
};
```

TRADITIONAL

```
sub rcx, rdx ; rcx = end-begin  
mov rax, rcx  
shr rax, 2    ; (end-begin)/4  
je .L4  
add rcx, rdx  
xor eax, eax
```

```
size_t size() const noexcept {  
    return _M_finish - _M_start;  
}
```

RANGE

```
xor eax, eax  
cmp rdx, rcx ; begin==end?  
je .L4
```

```
auto __begin = begin(v);  
auto __end = end(v);  
for (auto __it = __begin;  
     __it != __end;  
     ++it)
```

WALKTHROUGH

```
; rcx ≡ end, rdx = begin, eax = 0
.L3:
    add eax, DWORD PTR [rdx]      ; eax += *rdx
    add rdx, 4                    ; rdx += sizeof(int)
    cmp rdx, rcx                 ; is rdx == end?
    jne .L3                      ; if not, loop
    ret                          ; we're done
```

BACKSTORY

SO, WHICH APPROACH IS BEST?

ALSO

- Optimizer settings are important
- `std::accumulate`

WHAT HAS MY COMPILER DONE FOR ME LATELY?

MULTIPLICATION

```
int mulByY(int x, int y) {  
    return x * y;  
}
```

[View](#)

```
mulByY(int, int):  
    mov eax, edi  
    imul eax, esi  
    ret
```

MULTIPLICATION

```
int mulByConstant(int x) { return x * 2; }
```

[View](#)

MULTIPLICATION

```
int mulBy65599(int a) {  
    return (a << 16) + (a << 6) - a;  
    //           ^           ^  
    //   a * 65536      |  
    //                   a * 64  
    // 65536a + 64a - 1a = 65599a  
}
```

[View](#)

DIVISION

```
int divByY(int x, int y) {  
    return x / y;  
}  
int modByY(int x, int y) {  
    return x % y;  
}
```

[View](#)

Haswell 32-bit divide - 22-29 cycles!

```
divByY(int, int):  
    mov eax, edi  
    cdq  
    idiv esi  
    ret  
modByY(int, int):  
    mov eax, edi  
    cdq  
    idiv esi  
    mov eax, edx  
    ret
```

DIVISION

```
unsigned divByConstant(unsigned x) { return x / 2; }
```

[View](#)

DIVISION

```
mov eax, edi          ; eax = x
mov edx, 0xaaaaaaaaab
mul edx              ; eax:edx = x * 0xaaaaaaaaab
mov eax, edx          ; (x * 0xaaaaaaaaab) >> 32
                      ; ≡ (x * 0xaaaaaaaaab) / 0x10000000
                      ; ≡ x * 0.6666666667
shr eax              ; x * 0.3333333333
ret
```

MODULUS

```
int modBy3(unsigned x) {  
    return x % 3;  
}
```

[View](#)

```
mov eax, edi  
mov edx, 0xaaaaaaaaab  
mul edx  
mov eax, edx  
shr eax  
lea eax, [rdx+rdx*2]  
sub edi, eax  
mov eax, edi  
ret
```

WHY MODULUS?

- Bucket selection in hash maps
- libc++ special-cases power-of-two
- boost multi_index

COUNTING BITS

```
int countSetBits(int a) {  
    int count = 0;  
    while (a) {  
        count++;  
        a &= (a-1);  
    }  
    return count;  
}
```

[View](#)

SUMMATION

```
constexpr int sumTo(int x) {
    int sum = 0;
    for (int i = 0; i <= x; ++i)
        sum += i;
    return sum;
}
int main(int argc, const char *argv[]) {
    return sumTo(20);
}
```

[View](#)

SUM(X)

$$\begin{aligned}\sum_{n=0}^x n &\equiv \frac{x(x+1)}{2} \\ &\equiv x + \frac{x(x-1)}{2}\end{aligned}$$

WHAT ELSE HAS MY COMPILER DONE FOR ME LATELY?

FUNCTIONS

```
int someFunc(int);
int sum(const vector<int> &v) {
    int result = 0;
    for (size_t i = 0;
        i < v.size(); ++i)
        result += someFunc(v[i]);
    return result;
}
```

[View](#)

[Accumulate version](#)

```
.L3:
    mov edi, DWORD PTR [rdx+rbx*4]
    inc rbx
    call someFunc(int)
    mov rdx, QWORD PTR [rbp+0]
    add r12d, eax
    mov rax, QWORD PTR [rbp+8]
    sub rax, rdx
    sar rax, 2
    cmp rbx, rax
    jb .L3
```

ALIASING 1/2

```
// Worst API ever...
void sumToFirstZero(
    int *first, int &num,
    int &sum) {
    sum = 0;
    int i;
    for (i = 0; i < num; ++i) {
        if (first[i] == 0) break;
        sum += first[i];
    }
    num = i;
}
```

[View](#)

```
.L4:
    inc rax
    mov ecx, DWORD PTR [rdi-4+rax*4]
    test ecx, ecx
    je .L2
.L3:
    add r8d, ecx
    mov r9d, eax
    mov DWORD PTR [rdx], r8d
    cmp DWORD PTR [rsi], eax
    jg .L4
```

Alternative implementation

ALIASING 2/2

```
void maxArray(double* x, double* y)
    for (int i = 0; i < 65536; i++) {
        if (y[i] > x[i])
            x[i] = y[i];
    }
}
```

View
Better

```
maxArray(double*, double*):
    lea rax, [rsi+32]
    cmp rdi, rax
    jnb .L10
    lea rax, [rdi+32]
    cmp rsi, rax
    jb .L9
    ...
```

HEAP ELISION

```
int func() {  
    auto a = make_unique<int>(42);  
    auto b = make_unique<int>(24);  
    return *a + *b;  
}
```

[View](#)

TERNARY OPERATOR

```
int func(int i) {  
    return (i < 0) ? 1234 : 5678;  
}
```

[View](#)

```
func(int):  
    mov eax, edi  
    sar eax, 31  
    and eax, -4444  
    add eax, 5678  
    ret
```

VIRTUAL METHODS

```
struct Coster { virtual int costFor(int x) = 0; };
int totalCost(Coster &coster,
              const vector<int> &v) {
    int result = 0;
    for (auto i : v)
        result += coster.costFor(i);
    return result;
}
```

[View](#)

WHAT HAS MY COMPILER DONE FOR ME LATELY?

A lot!

HOW IT WORKS



HOW IT WORKS - BACKEND

- Written in node.js
- Runs on Amazon

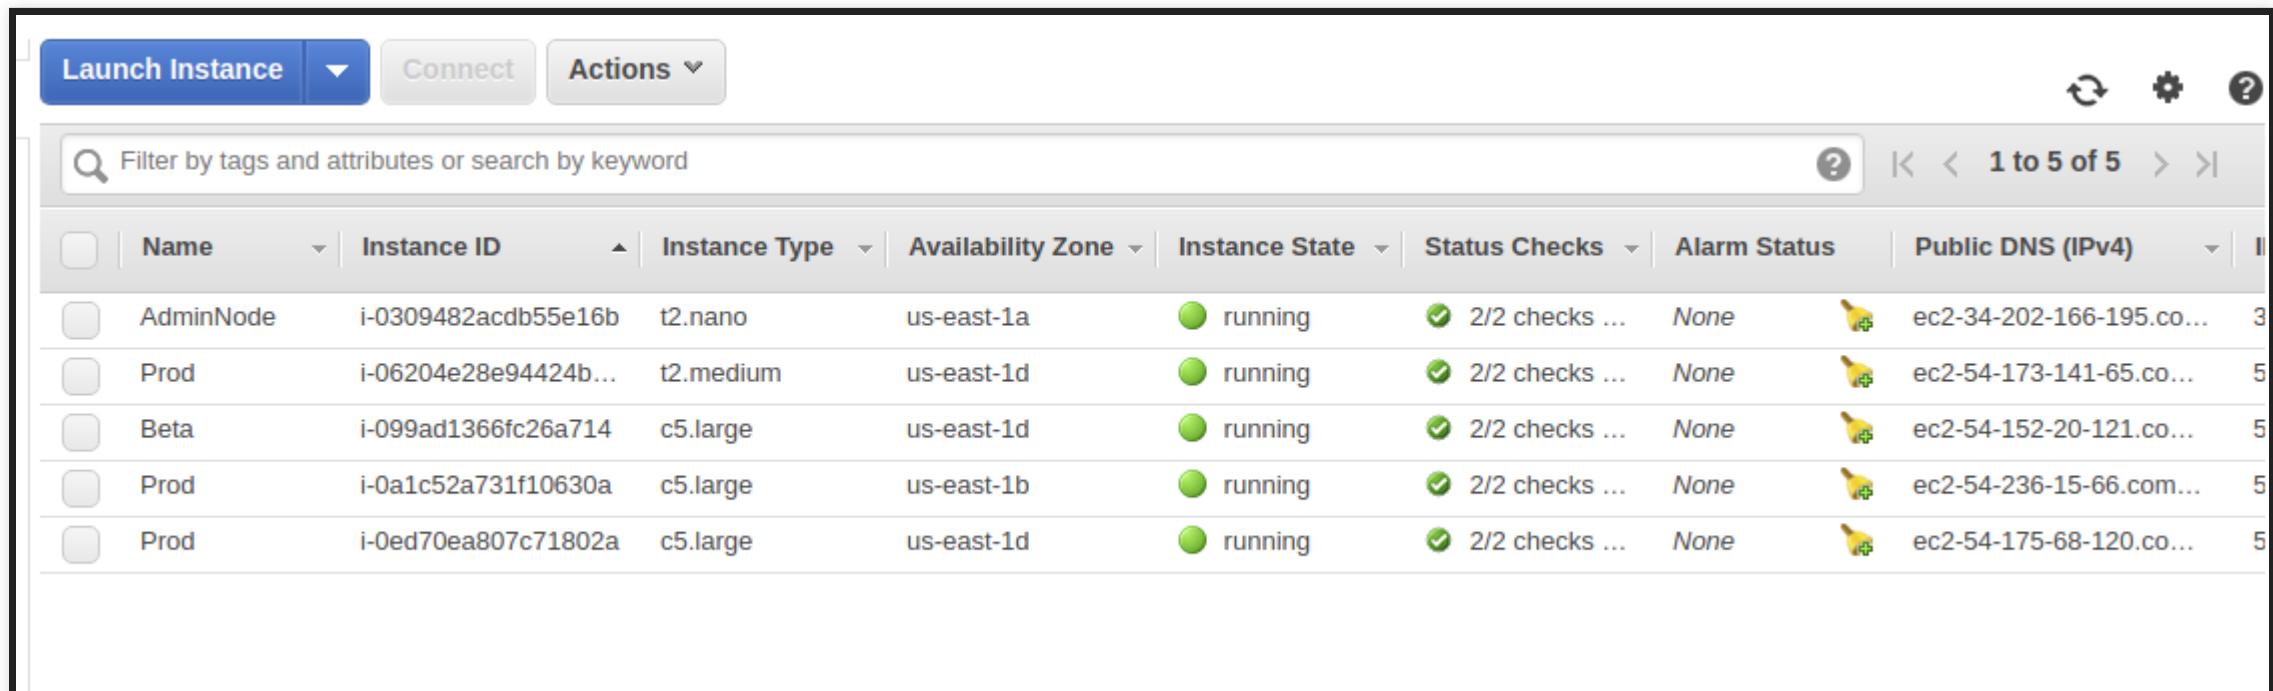
NODE.JS

```
function compile(req, res, next) {
  // exec compiler, feed it req.body, parse output
}
var webServer = express();
var apiHandler = express.Router();
apiHandler.param('compiler',
  function (req, res, next, compiler) {
    req.compiler = compiler;
    next();
});
apiHandler.post('/compiler/:compiler/compile', compile);
webServer.use('/api', apiHandler);
webServer.listen(10240);
```

AMAZON EC2

- Edge cache
- Load balancer
- Virtual machines
- Docker images
- Shared compiler storage

AMAZON EC2



The screenshot shows the AWS EC2 Instances page. At the top, there are buttons for "Launch Instance", "Connect", and "Actions". Below that is a search bar with the placeholder "Filter by tags and attributes or search by keyword". To the right of the search bar are navigation icons for "1 to 5 of 5" and other controls. The main area is a table listing five EC2 instances:

	Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS (IPv4)	
<input type="checkbox"/>	AdminNode	i-0309482acdb55e16b	t2.nano	us-east-1a	● running	✓ 2/2 checks ...	None	 ec2-34-202-166-195.co... 3	
<input type="checkbox"/>	Prod	i-06204e28e94424b...	t2.medium	us-east-1d	● running	✓ 2/2 checks ...	None	 ec2-54-173-141-65.co... 5	
<input type="checkbox"/>	Beta	i-099ad1366fc26a714	c5.large	us-east-1d	● running	✓ 2/2 checks ...	None	 ec2-54-152-20-121.co... 5	
<input type="checkbox"/>	Prod	i-0a1c52a731f10630a	c5.large	us-east-1b	● running	✓ 2/2 checks ...	None	 ec2-54-236-15-66.com... 5	
<input type="checkbox"/>	Prod	i-0ed70ea807c71802a	c5.large	us-east-1d	● running	✓ 2/2 checks ...	None	 ec2-54-175-68-120.co... 5	

THE COMPILERS

- Built through docker images
- Compilers stored on S3
- OSS ones publically available
- MS compilers via WINE

HOW IT WORKS - SECURITY

- Compilers: huge attack vector
- Principle of "what's the worst could happen"
- Docker
- LD_PRELOAD

HOW IT WORKS - FRONTEND

- Microsoft's Monaco
- GoldenLayout

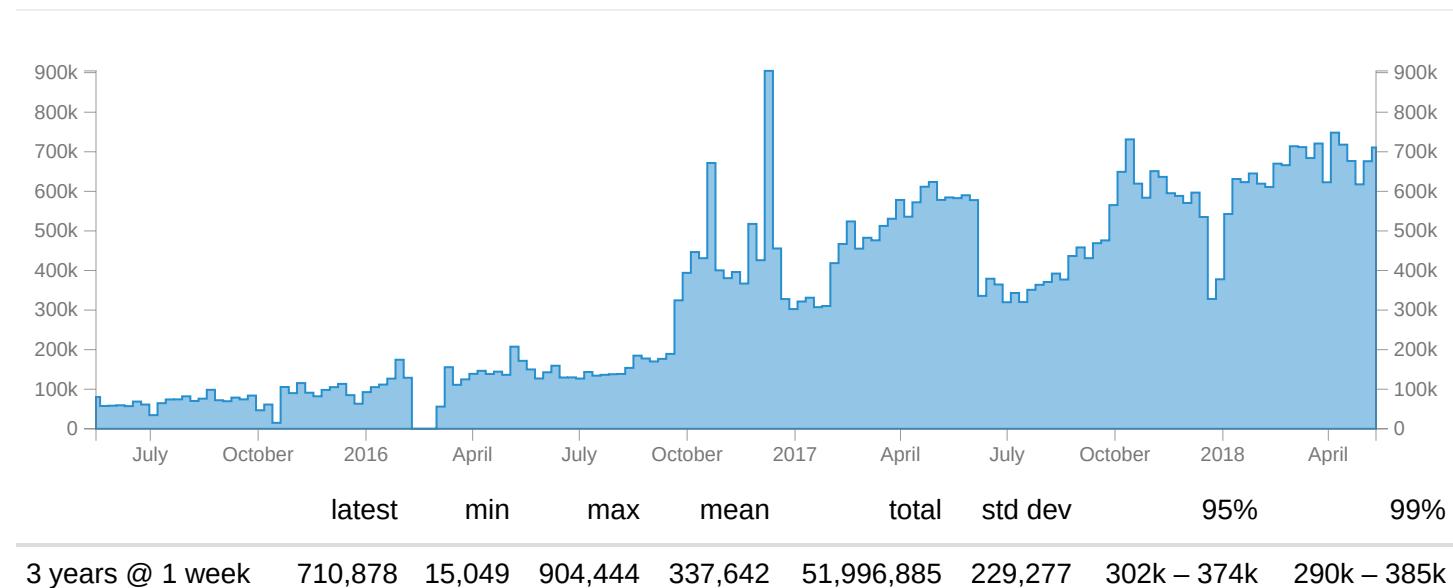
THE CODE

- github.com/mattgodbolt/compiler-explorer
- github.com/mattgodbolt/compiler-explorer-image
- Running locally is easy!

```
$ make
```

USAGE

All Compilations



StatHat tracks custom stats in your apps with one line of code.

- Up-to-the-minute charts
- Automatic anomaly detection
- 30-day forecasts

OTHER USES

- Code pastebin
- Compiler development
- C++ REPL
- Training resource

COMING SOON...

- Execution support

THANKS

- Thanks to contributors:
 - Rubén Rincón
 - Jared Wyles, Partouf, Simon Brand, Chedy Najjar, Filipe Cabecinhas, Johan Engelen...and the rest!
- Thanks to Patreon folks
- Thanks to awesome C++ community
- Thanks to you!

GO READ SOME ASSEMBLY!

godbolt.org

(AND THANK A COMPILER DEVELOPER)